

Task 2: Credit Risk Prediction

1. Introduction and Problem Statement

The goal of this project is to build a **classification model** that predicts whether a loan applicant is likely to **default** or **not**.

This is a critical task for financial institutions, as accurate credit risk prediction helps to **minimize financial risk** and ensures that loans are granted to individuals with a **high probability of repayment**.

```
# This Python 3 environment comes with many helpful analytics
# libraries installed
# It is defined by the kaggle/python Docker image:
# https://github.com/kaggle/docker-python
# For example, here's several helpful packages to load

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter)
# will list all files under the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/)
# that gets preserved as output when you create a version using "Save &
# Run All"
# You can also write temporary files to /kaggle/temp/, but they won't
# be saved outside of the current session

/kaggle/input/loan-data/loan_data.csv
```

2. Dataset Understanding and Description

We are using the **Loan Prediction Dataset**. This dataset contains information related to loan applicants, including the following categories:

- **Demographics:** Gender, Marital Status, Education, Dependents
- **Financials:** Applicant Income, Co-applicant Income, Loan Amount
- **Loan Details:** Loan Term, Credit History, Property Area

- **Target Variable:** Loan_Status (Y / N)

3. Data Cleaning and Preparation

We will start by loading the dataset and handling **missing values**. Since real-world data is often *messy*, missing **categorical variables** will be imputed using the **mode**, while missing **numerical variables** will be imputed using the **median**.

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report

# 1. Load the dataset
# Assuming the file is in the current directory as 'loan_data[1].csv'
df = pd.read_csv('/kaggle/input/loan-data/loan_data.csv')

# 2. Handle missing data appropriately
# Checking for null values
print("Missing values per column:\n", df.isnull().sum())

# If there were missing values, we would handle them here:
# df.fillna(df.median(numeric_only=True), inplace=True) # For numerical
# df.fillna(df.mode().iloc[0], inplace=True) # For categorical

# 3. Data Transformation: Convert 'purpose' column to dummy variables
# The 'purpose' column is categorical and needs to be numerical for the model
df_final = pd.get_dummies(df, columns=['purpose'], drop_first=True)

print("\nData Preparation Complete. Shape of dataset:",
df_final.shape)
df_final.head()

Missing values per column:
 credit.policy      0
 purpose            0
 int.rate           0
 installment        0
 log.annual.inc    0
 dti                0
 fico               0
 days.with.cr.line 0
 revol.bal          0
 revol.util         0

```

```
inq.last.6mths      0  
delinq.2yrs         0  
pub.rec             0  
not.fully.paid     0  
dtype: int64
```

Data Preparation Complete. Shape of dataset: (9578, 19)

```
    credit.policy  int.rate  installment  log.annual.inc  dti  
fico  \  
0           1   0.1189       829.10    11.350407  19.48   737  
1           1   0.1071       228.22    11.082143  14.29   707  
2           1   0.1357       366.86    10.373491  11.63   682  
3           1   0.1008       162.34    11.350407  8.10    712  
4           1   0.1426       102.92    11.299732  14.97   667
```

```
    days.with.cr.line  revol.bal  revol.util  inq.last.6mths  
delinq.2yrs  \  
0          5639.958333    28854      52.1          0  
0  
1          2760.000000    33623      76.7          0  
0  
2          4710.000000    3511       25.6          1  
0  
3          2699.958333    33667      73.2          1  
0  
4          4066.000000    4740       39.5          0  
1
```

```
    pub.rec  not.fully.paid  purpose_credit_card  
purpose_debt_consolidation  \  
0          0                  0          False  
True  
1          0                  0          True  
False  
2          0                  0          False  
True  
3          0                  0          False  
True  
4          0                  0          True  
False
```

```
    purpose_educational  purpose_home_improvement  
purpose_major_purchase  \  
0                      False          False
```

```

False
1      False      False
False
2      False      False
False
3      False      False
False
4      False      False
False

  purpose_small_business
0      False
1      False
2      False
3      False
4      False

```

4. Exploratory Data Analysis (EDA)

Visualization helps us understand the **relationship between features** and the **likelihood of loan approval**.

```

# Set visual style
sns.set(style="whitegrid")

# Create a figure for visualizations
plt.figure(figsize=(18, 5))

# Plot 1: Distribution of Installments (Key Feature: Loan Amount)
# We use 'installment' as the primary indicator of the loan
# size/burden
plt.subplot(1, 3, 1)
sns.histplot(data=df, x='installment', hue='not.fully.paid', bins=30,
kde=True, palette='viridis')
plt.title('Installment Distribution by Loan Status')

# Plot 2: Purpose of Loan (Key Feature: Education/Category)
# Note: 'educational' is one of the categories within the 'purpose'
# column
plt.subplot(1, 3, 2)
sns.countplot(data=df, y='purpose', hue='purpose', palette='magma',
legend=False)
plt.title('Frequency of Loan Purposes')

# Plot 3: Income (Key Feature: Log Annual Income)
# Fixing the FutureWarning by assigning x to hue
plt.subplot(1, 3, 3)
sns.boxplot(data=df, x='not.fully.paid', y='log.annual.inc',
hue='not.fully.paid', palette='Set2', legend=False)

```




```

plt.title('Log Annual Income vs. Loan Status')


plt.tight_layout()
plt.show()

# Additional Visualization: FICO Score vs. Interest Rate
# This helps visualize how risk (Credit Score) affects the cost of the
# loan
plt.figure(figsize=(10, 6))
sns.jointplot(x='fico', y='int.rate', data=df, color='purple',
kind='hex')
plt.suptitle('FICO Score vs. Interest Rate', y=1.02)
plt.show()

```



<Figure size 1000x600 with 0 Axes>



5. Model Training and Testing

We will split the data into **training and testing sets** (80/20) and train a **Logistic Regression** model, which is highly effective for **binary classification** tasks.

```
# Import necessary tools
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler

# 1. Defining Features (X) and Target (y)
X = df_final.drop('not.fully.paid', axis=1)
y = df_final['not.fully.paid']
```

```

# 2. Splitting the dataset (70% Train, 30% Test)
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.30, random_state=101)

# 3. Feature Scaling
# This step fixes the "ConvergenceWarning" by normalizing the data
range
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# 4. Initializing and Training the Logistic Regression Model
# With scaled data, the model converges much faster
model = LogisticRegression(max_iter=1000)
model.fit(X_train_scaled, y_train)

# 5. Making Predictions
predictions = model.predict(X_test_scaled)

print("Model training and testing phase completed successfully with
scaled data.")

Model training and testing phase completed successfully with scaled
data.

```

6. Evaluation Metrics

To assess how well our model performed, we check the **Accuracy** and examine the **Confusion Matrix** to evaluate the balance of **False Positives** and **False Negatives**.

```

# 1. Calculate Accuracy
accuracy = accuracy_score(y_test, predictions)
print(f"Overall Model Accuracy: {accuracy * 100:.2f}%")

# 2. Generate Classification Report (Precision, Recall, F1)
print("\nClassification Report:")
print(classification_report(y_test, predictions))

# 3. Visualizing the Confusion Matrix
cm = confusion_matrix(y_test, predictions)

plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Fully Paid (0)', 'Default (1)'],
            yticklabels=['Fully Paid (0)', 'Default (1)'])

plt.title('Confusion Matrix for Credit Risk Prediction')
plt.xlabel('Predicted Label')
plt.ylabel('Actual Label')
plt.show()


```

Overall Model Accuracy: 84.62%

Classification Report:

	precision	recall	f1-score	support
0	0.85	0.99	0.92	2431
1	0.52	0.04	0.07	443
accuracy			0.85	2874
macro avg	0.68	0.51	0.49	2874
weighted avg	0.80	0.85	0.79	2874

Confusion Matrix for Credit Risk Prediction



7. Conclusion: Summarizing Key Insights

- **Data Quality:** The dataset was clean with no missing values. Categorical features like purpose were successfully encoded using **one-hot encoding**.
- **EDA Insights:**

- Borrowers with higher FICO scores are significantly more likely to pay back their loans fully.
 - Certain loan purposes, such as **debt consolidation**, are more frequent but also show a noticeable volume of defaults.
 - Higher installments don't always correlate with higher default rates, as **income levels** also play a role.
- **Model Performance:** The **Logistic Regression** model achieved an accuracy of approximately **84.62%**. Examining the **confusion matrix** shows whether the model is better at predicting successful payments than defaults, which is common in **imbalanced datasets**.
 - **Recommendations:** For future improvements:
 - Handle **class imbalance** (fewer defaults than paid loans) using techniques like **SMOTE**.
 - Try alternative models such as **Random Forest Classifier** to enhance sensitivity to high-risk applicants.